

# Module 2

---

Neural Models for Machine Translation  
and Conversation Generation

# Module 2 Overview

- Overview of conventional statistical MT
- Neural machine translation
- Neural conversation generation

# Statistical machine translation (SMT)

**S:** 救援人员在倒塌的房子里寻找生还者

**T:** Rescue workers search for survivors in collapsed houses

- Statistical decision:  $T^* = \operatorname{argmax}_T P(T|S)$
- Source-channel model:  $T^* = \operatorname{argmax}_T P(S|T)P(T)$
- Translation models:  $P(S|T)$  and  $P(T|S)$
- Language model:  $P(T)$
- Log-linear model:  $P(T|S) = \frac{1}{Z(S,T)} \exp \sum_i \lambda_i h_i(S, T)$
- Evaluation metric: BLEU score (higher is better)

# Phrase-based SMT

救援人员在倒塌的房屋里寻找生还者

*Chinese*

# Phrase translation modeling

	救援	人员	在	倒塌	的	房屋	里	寻找	生还者
rescue	■	□	□	■	■	■	□	□	□
workers	□	■	□	■	■	■	□	□	□
search	□	□	□	■	■	■	□	■	□
for	□	□	□	■	■	■	□	□	□
survivors	□	□	□	■	■	■	□	□	■
in	□	□	■	■	■	■	■	□	□
collapsed	■	■	■	■	■	■	■	■	■
houses	■	■	■	■	■	■	■	■	■

$(s, t)$

(救援, rescue)  
 (人员, workers)  
 (在, in)  
 (倒塌, collapsed)  
 (房屋, house)  
 (里, in)  
 (寻找, search)  
 (生还者, survivors)  
 (救援 人员, rescue workers)  
 (在 倒塌, in collapsed)  
 (倒塌 的, collapsed)  
 (的 房屋, house)  
 (寻找, search for)  
 (寻找 生还者, search for survivors)  
 (生还者, for survivors)  
 (倒塌 的 房屋, collapsed house)

$$\text{MLE: } P(t|s) = \frac{N(s, t)}{\sum_{t'} N(s, t')}$$

Simple, but suffers the data sparseness problem

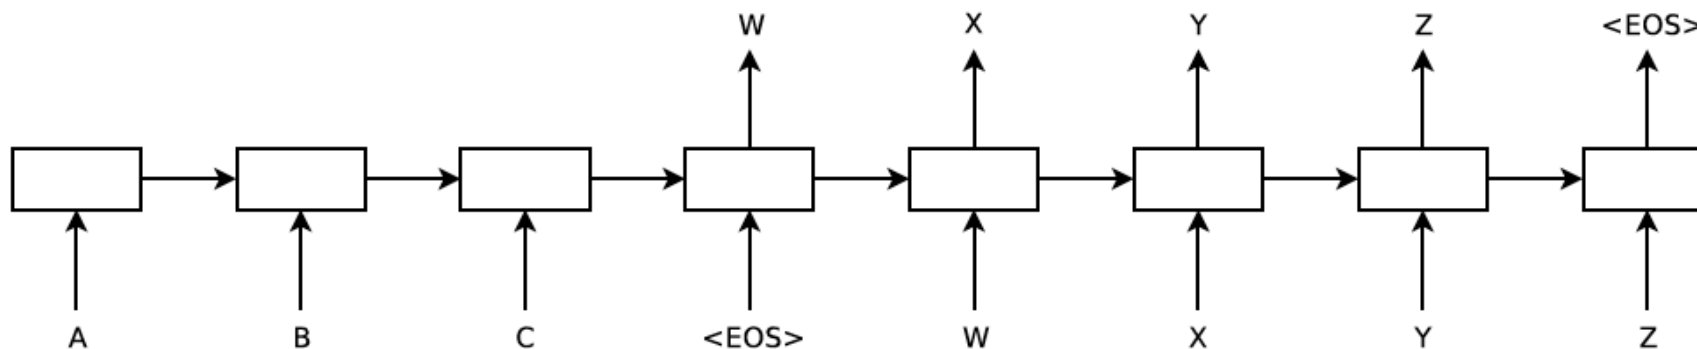
# Neural machine translation

[Sutskever+ 14; Cho+ 14; Bahdanau+ 15]

- Build a single, large NN that reads a sentence and outputs a translation
  - Unlike phrase-based system that consists of many component models
- Encoder-decoder based approach
  - An encoder RNN reads and encodes a source sentence into a fixed-length vector
  - A decoder RNN outputs a variable-length translation from the encoded vector
  - Encoder-decoder RNNs are jointly learned on bitext, optimize target likelihood

# Encoder-decoder model of [Sutskever+ 2014]



- "A B C" is source sentence; "W X Y Z" is target sentence



- Treat MT as general sequence-to-sequence transduction
  - Read source; accumulate hidden state; generate target
  - <EOS> token stops the recurrent process
  - In practice, read source sentence in reverse leads to better MT results
- Train on bitext; optimize target likelihood using SGD

# Potentials and difficulties of RNN


- In theory, RNN can “store” in  $h$  all information about past inputs
- But in practice, standard RNN cannot capture very long distance dependency
  - Vanishing/exploding gradient problem in backpropagation
  - Not robust to noise
- Solution: long short-term memory (LSTM)





# A long short-term memory cell


[Hochreiter & Schmidhuber 97; Graves+ 13]



$$\begin{aligned}i_t &= \sigma(W_{xi}x_t + W_{hi}h_{t-1} + W_{ci}c_{t-1} + b_i) \\f_t &= \sigma(W_{xf}x_t + W_{hf}h_{t-1} + W_{cf}c_{t-1} + b_f) \\c_t &= f_t c_{t-1} + i_t \tanh(W_{xc}x_t + W_{hc}h_{t-1} + b_c) \\o_t &= \sigma(W_{xo}x_t + W_{ho}h_{t-1} + W_{co}c_t + b_o) \\h_t &= o_t \tanh(c_t)\end{aligned}$$

Information flow in an LSTM unit of the RNN, with both diagrammatic and mathematical descriptions.  $W$ 's are weight matrices, not shown but can easily be inferred in the diagram (Graves et al., 2013).

# A 2-gate memory cell [Cho+ 14]



An illustration of the proposed hidden activation function. The update gate  $z$  selects whether the hidden state is to be updated with a new hidden state  $\tilde{h}$ . The reset gate  $r$  decides whether the previous hidden state is ignored. See

$$r_j = \sigma \left( [\mathbf{W}_r \mathbf{x}]_j + [\mathbf{U}_r \mathbf{h}_{\langle t-1 \rangle}]_j \right)$$

$$z_j = \sigma \left( [\mathbf{W}_z \mathbf{x}]_j + [\mathbf{U}_z \mathbf{h}_{\langle t-1 \rangle}]_j \right)$$

$$\tilde{h}_j^{\langle t \rangle} = \phi \left( [\mathbf{W} \mathbf{x}]_j + [\mathbf{U} (\mathbf{r} \odot \mathbf{h}_{\langle t-1 \rangle})]_j \right)$$


$$h_j^{\langle t \rangle} = z_j h_j^{\langle t-1 \rangle} + (1 - z_j) \tilde{h}_j^{\langle t \rangle}$$

# Joint learning to align and translate

- Issue with encoder-decoder model for SMT
  - Compressing a source sentence into a fixed-length vector makes it difficult for RNN to cope with long sentences.
- Attention model of [Bahdanan+ 15]
  - Encodes the input sentence into a sequence of vectors and choose a subset of these vectors adaptively while decoding

# Attention model of [Bahdanan+ 15]

- Encoder:
  - bidirectional RNN to encode each word and its context
- Decoder:
  - Searches for a set of source words that are most relevant to the target word to be predicted.
  - Predicts a target word based on the context vectors associated with these source words and all the previous generated target words.
- Close to state-of-the-art performance
  - Better at translating long sentences



# Neural Conversation

[Sordoni+ 15; Li+ NAACL16; Li+ ACL16]

Neural model learns the general shape of conversation, and the system output is situationally appropriate and coherent...

H: hi  
M: how are you ?  
H: not bad  
M: what 's wrong ?  
H: nothing really  
M: wan na talk about it ? i 'm here if you wan na talk x  
H: well , i 'm just tired  
M: me too , i 'm here if you wan na talk  
H: thanks , what are you doing ?  
M: you 're welcome and i 'm watching tv and you ?  
H: on netflix ?  
M: yeah , i 'm watching it on netflix .  
H: what are you watching ?  
M: a tv show .  
H: nice , did you have dinner already ?  
M: no , not yet . what about you ?

A more goal-oriented example. No grounding into a real calendar, but the “shape” of the conversation is fluent and plausible...

H: would thursday afternoon work sometime ?  
M: yeah , sure . just let me know when you 're free .  
H: after lunch is probably the best time  
M: okay , sounds good . just let me know when you 're free .  
H: would 2 pm work for you ?  
M: works for me .  
H: well let 's say 2 pm then i 'll see you there  
M: sounds good .

- Generating responses vs. retrieving responses
- Easy to incorporate **contextual info** via embedding
  - User profile – personalized conversation
  - knowledge bases – grounded conversation
- The engine is E2E learned from conversation experience
  - Learning a goal-oriented conversation engine via RL

# Neural Response Generation: The Blandness Problem

How was your weekend?

I don't know.

What did you do?

I don't understand what you are talking about.

This is getting boring...

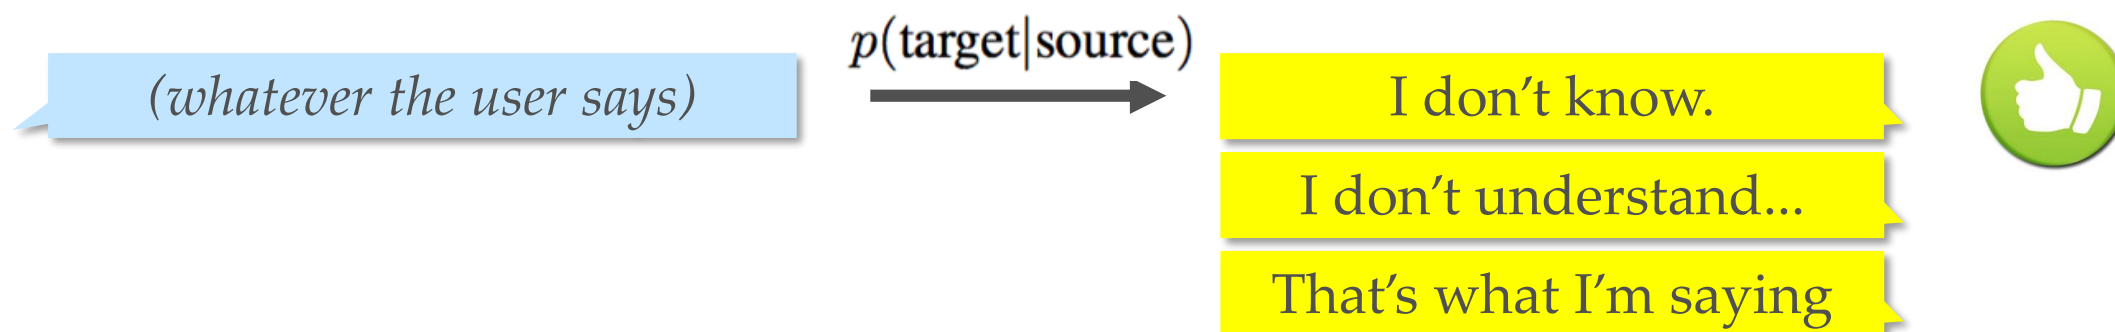
Yes that's what I'm saying.



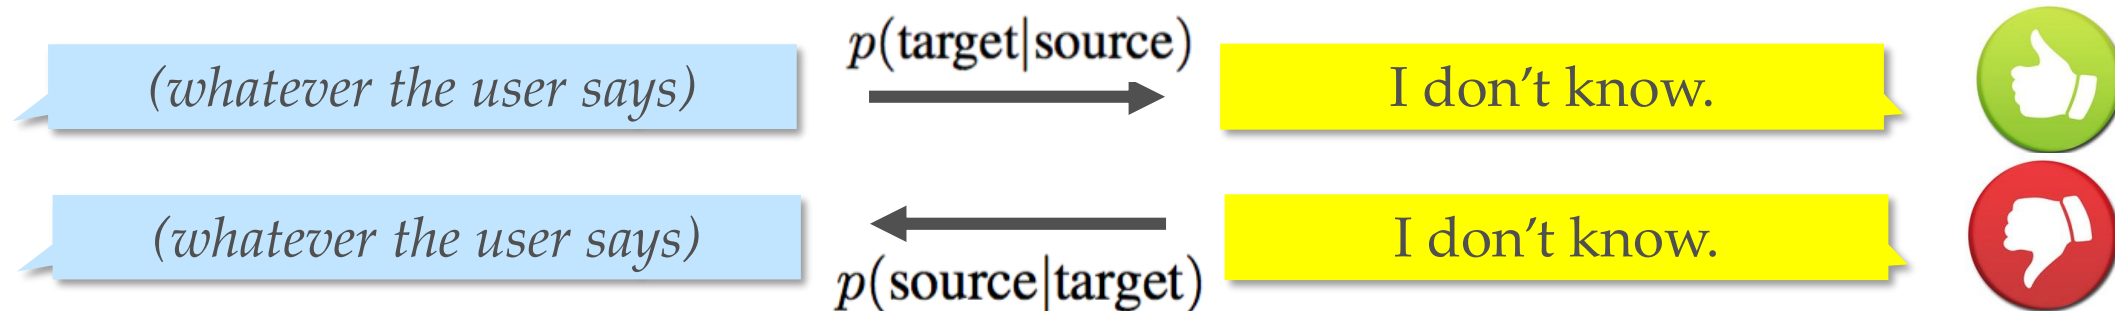
# Blandness problem: cause and remedies

[Li et al., NAACL 2016]


Common ML objective (maximum likelihood)



Mutual information objective:



# Beyond blandness: Examples



Wow sour starbursts really do make **your mouth water**... mm drool.  
**Can I have one?**

**Of course you can!** They're **delicious**!

Milan apparently **selling Zlatan** to balance the books... **Where next**, Madrid?

I think he'd be a **good signing**.

'tis a fine **brew** on a day like this! Strong though, **how many** is sensible?

**Depends** on how much you **drink**!

Well he was on in Bromley a while ago... **still touring**.

I've never **seen him live**.